

S3 Enhancement Delivery

Test & walkthrough guide — all scenarios

Generated 3 August 2026 · repo `ams-s3-demo` · branch `s3-console-stage-routes`
Companion docs: `docs/design/S3_DESIGN.md` (technical design) · `demo/DEMO_TEST_GUIDE.md` (presenter script) · `repos/README.md` (the target repositories, and onboarding a new one)

Environment

Six processes make up the full environment. The console and the app a given user story targets must be running; the rest can stay down if you are not showing that user story.

PORT	SERVICE	URL	START COMMAND (TERMINAL BY TERMINAL)
8000	FastAPI backend the engine — never on screen	localhost:8000	<code>apps/run-console.sh</code> starts 8000 and 5173 together
5173	AMS Console — start here login, then <code>/s3</code> ; managers also get <code>/admin</code>	localhost:5173	<code>apps/run-console.sh</code>
8501	PolicyCore portal (user story-041 / user story-042 target) "the client's app"	localhost:8501/ <code>sl_policycore</code> base path — the bare root 404s	<code>apps/run-policycore.sh</code>
8081	ClaimsPortal Contracts Team console (user story-043)	localhost:8081	<code>apps/run-policy-service.sh</code>
8082	ClaimsPortal Claims Team console (user story-043)	localhost:8082	<code>apps/run-claims-service.sh</code>
8083	EnrolDirect (user story-045 target) online enrolment channel	localhost:8083	<code>apps/run-enroldirect.sh</code>

Both ClaimsPortal services together, in one foreground terminal: `demo/run_s3_claimsportal.sh`. It traps `kill 0` on exit, so use the two separate commands above if you want them in independent background terminals. `demo/run_mockapp.sh` is an equivalent launcher for the portal — same base path, same `PYTHONPATH` export — kept because teammates and the admin panel both invoke it by name.

YOU NO LONGER NEED A TERMINAL FOR MOST OF THIS

Log in as **Manager** and open `/admin`. It shows which of the four target apps are answering on their port, starts or stops them, clears logs, runs the reset scopes, and onboards a repo by writing its `.s3targets.json`. Service status is a plain TCP connect — no `ps`, no `lsof` — so it works on a locked-down host.

Two limits worth knowing before you rely on it mid-run. There is **no service id for the console**: it cannot restart the process serving your request, and naming it is a 400 rather than an attempt. And a manifest the panel just wrote **needs a console restart** before that target registers.

Where the code lives

Two folders, and the split is load-bearing. `repos/` holds the three target repositories — `policycore`, `claimsportal`, `enroldirect` — which are what S3 *changes*. `apps/` holds the console and the launch scripts, which are what does the changing. Dropping a directory into `repos/` with a `.s3targets.json` manifest registers it as a target at next start, with no edit to any registry file; putting its user story under `stories/` opens the board ticket. See `repos/README.md` for the manifest contract and the four onboarding steps — it is short, and it is the authority.

Direct URLs for showing raw data

URL	SHOWS
localhost:8000/docs	OpenAPI browser for the whole S3 API
localhost:8081/api/policies	All ClaimsPortal policies as JSON
localhost:8081/api/policies/MS-1004	One policy — the clearest place to show <code>deductible</code> appearing
localhost:8082/api/claims	Submitted claims with decision and <code>payableAmount</code>
localhost:8083/api/analysis/prospect-impact	EnrolDirect's both-policies comparison — the analysis user story-045 acts on

THREE THINGS THAT WILL BITE YOU

- Never run `uvicorn` with `--reload`.** The pipeline writes `.py` files into the tree `uvicorn` watches; the reload wipes the in-memory session store and every request then 401s while the UI still looks logged in.
- The ClaimsPortal processes serve whatever was last running.** There's no build step (plain Python), but `uvicorn` is deliberately started without `--reload`, so an Apply changing the source does not take effect until you restart both services. A process already running survives the reset entirely: it keeps serving the *post-user story* app while the source sits at baseline, and both ports answer 200 the whole time. Port-is-up is not process-is-current — check the process start time against your reset time.
- Start the portal via a launch script, never `streamlit` run by hand.** `apps/run-policycore.sh` and `demo/run_mockapp.sh` both export `PYTHONPATH="$PWD"` first and pass the same base path; either is fine, neither is optional. Streamlit otherwise puts `repos/policycore/` on `sys.path` instead of the repo root, and the app dies at `ModuleNotFoundError: No module named 'common'`. This one hides: a portal already running from a bad launch keeps serving its last good render until something re-executes the script — and a reset rewriting `app.py` is exactly that, so the error surfaces *during the reset* and looks like the reset broke it. It didn't.

Logins

Fictional roster; passcode is `1001 + roster position`.

NAME	PASSCODE	ROLE IN THESE TESTS
Ravi Kumar	1001	Developer — owns AMS-102, AMS-103
Elena Cruz	1002	Second engineer — owns AMS-098
Priya Nair	1003	Tester — owns AMS-101; QA hand-off target
Tom Becker	1004	Second tester

Seeded board

Verified against a live `GET /api/s3/jira/board` on 3 August 2026.

TICKET	USER STORY	STATUS	ASSIGNEE	TARGET	ORIGIN
AMS-101	US-2026-041	QA	Priya Nair	polycycore — plan tier upgrade	seeded
AMS-102	US-2026-042	In Progress	Ravi Kumar	polycycore — amendment priority	seeded
AMS-103	US-2026-043	To Do	Ravi Kumar	claimsportal — claims deductible	seeded
AMS-104	US-2026-044	To Do	Ravi Kumar	polycycore — flag urgent amendments	seeded
AMS-098	—	Done	Elena Cruz	none (ad-hoc path)	seeded
AMS-1045	US-2026-045	To Do	unassigned	enroldirect — prospect access	picked up from stories/

AMS-1045 IS NOT SEEDED — IT IS DERIVED, AND THAT IS A BEAT

A `.md` file dropped into `stories/` opens a board ticket on its own. The key is a pure function of the user story identifier (`US-2026-045` → **AMS-1045**), so it survives restarts and resets with nothing to persist, and it sits in a band the hand-seeded tickets cannot reach. It lands **unassigned** on purpose — that is what puts it in front of the manager, who sees exactly the unassigned ones with an Assign control.

Assignment used to be set-once, enforced only by hiding the button. A manager can now reassign or unassign, and `POST /s3/jira/assign-ticket` is **manager-only server-side** (`require_manager`) — it previously had no role check at all. If you rehearse the assign flow, the board you hand to the next presenter will show whoever you picked, not *unassigned*; reset the `tickets` scope to put it back.

Reset between runs

The pipeline mutates real files. Stop the services first, then run these in **this order** — the user story-042 reset must come *after* `reset_s3.sh`, because it restores the superset PolicyCore state that leaves both AMS-101 and AMS-102 ready to run.

```
demo/reset_s3.sh           # user story-041 + shared state: .cache/llm, out/, events, Jira caches
demo/reset_s3_endorsement.sh # user story-042 baseline (MUST run after reset_s3.sh)
demo/reset_s3_claimsportal.sh # user story-043 target, from repos/claimsportal/.baseline/
demo/reset_s3_enroldirect.sh # user story-045 target, from repos/enroldirect/.baseline/
```

`reset_s3_endorsement.sh` keeps its pre-reskin filename deliberately — teammates invoke it by name, and only its contents changed. The same reset scopes are on the manager's `/admin` panel if you would rather not leave the browser; there is no "reset everything" scope, by design — you name one.

Then restart the services. Hard-refresh the browser tab — the console caches per-ticket state in `localStorage` and drops it when the reset marker changes.

THE TWO POLICYCORE RESETS RESTORE FROM HEAD

demo/reset_s3.sh and demo/reset_s3_endorsement.sh restore source with `git checkout HEAD -- repos/...`, so **HEAD must already carry the paths they name**. Move a target directory and both stop at the checkout with *"pathspec did not match"* until the move is committed — committing it is the whole fix, nothing in the scripts needs changing.

The admin panel checks that before it runs anything and reports a named `reset_blocked_reason` rather than surfacing a raw git error out of a button. The ClaimsPortal and EnrolDirect resets copy from their committed `.baseline/` snapshots, so they never depend on HEAD.

EXPECTED AFTER A RESET

`git status` shows `repos/policycore/{app,core/db,core/amendments,core/models}.py` as modified. That is **correct** — repo HEAD is committed in a post-US-2026-042 state, so the pre-user story baseline necessarily differs from it.

Resets also clear `.cache/llm`, so impact analysis, design doc and release notes make live provider calls again (~10–20 s each on first click). Codegen and testgen always replay — instant and free. Run `demo/warm_s3_cache.sh` if you want the narrative beats instant too.

Before → after: what changes on screen

The "show the client their own app changing" moment for each scenario. Have this open while presenting.

SCENARIO	WHERE TO LOOK	BEFORE — BASELINE	AFTER APPLY	TO SEE IT
AMS-101 US-2026-041	Portal :8501/sl_policycore Policy Detail	No plan tier anywhere in the app.	Tier per contract (Standard / Premium / <i>the name the audience picked</i>), upgrade control on the detail view, monthly contribution recalculates and persists.	Browser refresh
AMS-102 US-2026-042	Portal :8501/sl_policycore Request a Contract Amendment	Exactly 5 fields — amendment type, describe the requested change, effective date, contact phone, contact email.	A 6th field, Priority (Standard / Urgent, defaults Standard). Existing submit flow unchanged.	Browser refresh
AMS-103 US-2026-043	Claims :8082 and Policy :8081	Policies carry no deductible. An \$80 claim on MS-1004 → ACCEPTED .	Policies gain a deductible (MS-1001 500, MS-1002 1000, MS-1003 500, MS-1004 100). The same \$80 claim on MS-1004 → REJECTED_BELOW_DEDUCTIBLE . A \$1,200 claim on MS-1001 → ACCEPTED, payableAmount 700 .	Restart both services

AMS-1045 US-2026-045	EnrolDirect :8083 Access check	A prospect is refused, and the refusal is an omission rather than a rule: "category PROSPECT has no online enrolment preference", requiredPreference: null. 0 of 6 prospects admitted.	A prospect resolves through the configured policy to the Guest preference. AP-4003 and AP-4008 are admitted with authorisingPreference: "Online Enrolment – Guest"; the other four are still refused, each for its own named reason (member-only sponsor, no preferences, lapsed contract). 2 of 6 admitted.	Restart the service
-------------------------	-----------------------------------	---	--	----------------------------

WHY USER STORY-045 IS THE DIFFERENT ONE

The other three user stories add something. **US-2026-045's baseline is a removal** — the checked-in source is the state after the impact analysis and before anyone acted on it, so the "before" is a refusal nobody chose. The user story settles a classification rather than building a feature, which is why the honest before/after is a decision changing, not a control appearing.

Second thing worth saying out loud: `repos/enroldirect/impact.py` is in the target's **core files but deliberately not in its codegen allowlist**. The model must read the analysis to understand the change and must not edit it — the analysis has to keep sizing *both* options after one is adopted. Core-file recall gets it into the prompt; the allowlist keeps it out of the diff, and this target has its own validator that fails loudly if a read-only file comes back modified.

THE EASIEST WAY TO MAKE A WORKING RUN LOOK BROKEN

The two PolicyCore scenarios need only a browser refresh — Streamlit re-executes `app.py` on each run.

Scenario 3 does not. The running processes serve whatever was loaded at startup (no `--reload`), so a just-applied user story is invisible on `:8081:8082` until you restart. Apply the change, switch to `:8082`, and nothing has changed. Ctrl-C both terminals, then `demo/run_s3_claimsportal.sh`.

Confirm the baseline is genuinely "before"

Run from the repo root after the resets, before you present. Every line must match the comment. Re-derived and re-run against the GRS-reskinned tree on 3 August 2026 — the old `coverage_tier / COVERAGE_TIERS` `grep` now passes vacuously, because those symbols no longer exist under any state of the repo. The user story-041 sentinel is `PLAN_TIERS / upgrade_tier`, not `plan_tier`: the field exists on `Policy` before the user story and the portal's contract header displays it, so `plan_tier` alone is a false positive. What user story-041 adds is the tier *ladder* and the upgrade call.

```
grep -c 'PLAN_TIERS\|upgrade_tier' repos/policycore/app.py # 0      - no user story-041 yet
grep -c 'selectbox("Priority")' repos/policycore/app.py    # 0      - no user story-042 yet
ls repos/policycore/core/tiers.py                          # absent - user story-041 creates it
ls repos/claimsportal/claims_service/claim_rules.py        # absent - user story-043 creates it
ls repos/policycore/core/amendments.py                     # PRESENT - user story-042's "before" needs it
grep -c 'DEFAULT_PROSPECT_POLICY' repos/enroldirect/applicants.py # 0      - no user story-045 yet
curl -s localhost:8081/api/policies/MS-1004                # no "deductible" key
curl -s -X POST localhost:8083/api/eligibility/check \
  -H 'Content-Type: application/json' -d '{"applicantId":"AP-4003"}' # "granted": false
ls s3_enhancement/out/ | wc -l                             # 0
```

Two of those are worth reading rather than glancing at. The user story-041 `grep` is scoped to **`app.py` only**, and its sentinel is `PLAN_TIERS / upgrade_tier` rather than `plan_tier`, for the same reason:

`core/models.py` and `core/db.py` carry a `plan_tier` column at baseline by design, so a `plan_tier` grep is not zero and is not a failure. The two ladder symbols appear in no source file at baseline — only in `AGENTS.md/CLAUDE.md`, which the greps do not read. And the `AP-4003` check must come back `"granted": false` with `"requiredPreference": null`; a prospect refused for any *other* reason means the EnrolDirect reset did not take.

Then confirm the services are up *and fresh* — every start time must be later than your last reset:

```
for p in 8000 5173 8501 8081 8082 8083; do
  url="http://localhost:$p/"
  [ "$p" = 8501 ] && url="http://localhost:8501/sl_policycore/" # Streamlit base path
  pid=$(lsof -ti tcp:$p -sTCP:LISTEN | head -1)
  code=$(curl -s -o /dev/null -w '%{http_code}' --max-time 5 "$url")
  printf "%-5s http=%-3s pid=%-6s %s\n" "$p" "$code" "${pid:-DOWN}" \
    "$(ps -o lstart= -p ${pid:-0} 2>/dev/null)"
done
```

Use `-sTCP:LISTEN`. Without it a Chrome tab connected to :5173 matches before the real listener and you read the wrong process's start time. The admin panel's service list answers the up/down half of this without a terminal, but it is a TCP probe only — it cannot tell you a process is *stale*, which is the failure this loop exists to catch.

Scenarios

THE SCREEN CHANGED — ARTIFACTS NOW OPEN IN A POP-UP

Acting on the client's "there is a lot of data on the screen... open it in a pop-up" feedback, every produced **artifact** now opens in a modal rather than rendering inline: release notes, deployment plan, design doc, test scenario table, per-case checklists, mutation diff, traceability matrix. What stays in the main flow is what you steer with — action buttons, verdict lines, token-cost lines. The left-hand stage rail is unchanged.

Practical effect on these scripts: where a step below says "read the X", expect to click through and then close the modal before the next button is reachable. Nothing moved stage to stage.

01 AMS-103 — ClaimsPortal, full pipeline

Start as: Ravi Kumar (1001) → hand off to Priya Nair (1003) · **Time:** ~15 min · **Best single end-to-end run**

1. Open :5173 → *Open application* → click **AMS-103**.
 2. **Run AI impact analysis**. It asks up to two clarifying questions — a gap question, then an assumption question. Answer each in the same box.
 3. Close the modal → **Stage 1 · Generate the change**.
 4. **Apply to repo**.
 5. **Stage 2 · Draft design doc** → **Assign tester & move to QA** (Priya Nair).
 6. Log out → log in as **Priya Nair** → open AMS-103 → **Stage 3 · Generate tests** → **Run tests** (real pytest, the same runner every target uses).
 7. **Prove the tests catch bugs** → **Stage 4 · Draft release notes** → mark Done.
- ❑ After the second answer it *must* produce a final analysis — the two-turn cap is hard.
 - ❑ Activity tab logs: clarification requested x2 → impact analysis drafted → status In Progress.
 - ❑ Selection panel reads **Files in this app 5 · used 5** and *"No subsystem doc matched closely enough"* — no `repos/policycore/systems/...` entry.
 - ❑ Token panel says *"no saving over whole-app context here"* — the honest branch.
 - ❑ Six Python files, each with a one-line reason.
 - ❑ Post-apply link reads **"open the Claims Team console"** → :8082 (not the PolicyCore portal).
 - ❑ On disk: `claim_rules.py` created, `deductible` present in `policy.py`.
 - ❑ Stage 3 was *locked* for Ravi — the developer cannot verify their own change.
 - ❑ Per-case checklist parsed from pytest's JUnit XML output, all green.
 - ❑ Mutation flips `<=` to `<`; the "claim at exactly the deductible" test goes red; file auto-restored (`git diff` on `claim_rules.py` is clean afterwards).

:8081 WILL NOT CHANGE — BY DESIGN

US-2026-043 line 50 states *"Both team consoles must keep rendering without changes to their HTML"*, and no `.html` is in the target's codegen allowlist. The policy console's table has no Deductible column, so it looks identical after the user story. The change is there — `GET :8081/api/policies` now returns `"deductible":500`.

The visible proof is on :8082. Restart the services, then submit a claim against MS-1001 (limit 25 000, deductible 500):

CLAIM AMOUNT	BEFORE THE USER STORY	AFTER THE USER STORY
400	ACCEPTED	REJECTED_BELOW_DEDUCTIBLE
500 — exactly the deductible	ACCEPTED	REJECTED_BELOW_DEDUCTIBLE
1 200	ACCEPTED	ACCEPTED — payable 700 now recorded
99 000	REJECTED_OVER_LIMIT	REJECTED_OVER_LIMIT

02 AMS-101 — the scoring showcase

As: Priya Nair (1003) — the ticket is seeded to her, already in QA, so she runs the whole flow · **Time:** ~10 min

1. Click **AMS-101** → **Run AI impact analysis**.
 2. **Stage 1 · Generate** — type an audience-chosen tier name (≤ 24 chars, no quotes or slashes, e.g. *Platinum*).
 3. **Apply** → post-apply runs `python -m repos.policycore.core.seed`.
 4. **Design doc** → **Generate tests** → **Run** (pytest) → **Mutation check** → **Release notes**.
- ☐ Panel reads **Files in this app 58 · used 4**.
 - ☐ Expand "7 other subsystems screened out" → bar chart, struck through: settlement 0.35, enrolment 0.22, risk 0.21, billing 0.20, underwriting 0.19, audit 0.15, reporting 0.12 — all under the 0.45 floor. **52 files never opened** — the 50 Java decoys plus the two Python modules of PolicyCore's own *enrolment* subsystem, which is the better line: screening is a scope decision, not a file-extension filter.
 - ☐ Token panel shows a real multiplier against whole-app context (3.4× on this user story — 18,310 estimated file tokens down to 5,336).
 - ☐ Your chosen tier name appears verbatim in the generated `tiers.py` — one recording serves any name.
 - ☐ Refresh **:8501/sl_policycore** → the portal now has the plan-tier upgrade control, and the monthly contribution moves with the tier.
 - ☐ Mutation weakens the same-tier guard `<=` → `<`; the suite goes red, then reverts.

03 AMS-102 — amendment priority

As: Ravi Kumar (1001) → hand off to a tester · Same step sequence as Scenario 01

THE DOCSTRING PROBLEM IS FIXED — THIS IS NOW SAFE TO APPLY ON STAGE

The old warning here said the user story-042 recording stripped roughly six function docstrings from `repos/policycore/core/db.py`, removing 81 lines to add 12. The 3 August 2026 re-record fixed it. Re-verified by diffing the committed recording against the baseline tree: **docstring counts are identical file for file** (`db.py` 40 before and after), and the whole four-file diff is **+16 / -32** rather than **+12 / -81**.

What is left is cosmetic and worth naming before a reviewer spots it: the model collapses PEP 8's two blank lines between top-level definitions down to one, and it de-indents the continuation line of two docstrings (`core/amendments.py`, `app.py`). Annoying in a diff; not a behaviour change. Model docstring mangling recurs on every re-record — always diff the recording against the baseline after one.

- ☐ All four files show pending changes. (Earlier recordings dropped one as byte-identical; this one does not.)
- ☐ `models.py` gains `priority: str = "Standard"`, `db.py` gains the column and the `INSERT` parameter, `amendments.py` takes it as a keyword argument.
- ☐ If applied: **:8501/sl_policycore** → *Request a Contract Amendment* gains a *Priority* dropdown (Standard / Urgent, Standard default).

04 The review loop

Where: any ticket, Stage 1, after Generate and before Apply

ACTION	EXPECTED
Show diff on a file	Per-file unified diff, collapsible
Ask about this file → <i>"why did you delete those blank lines?"</i>	Answers from its own diff — admits the deletion instead of denying it. (On user story-042 that is a real, visible change to point at.)
Ask → <i>"rename that variable to X"</i>	Returns a revised file; still staged, still not applied
Ask a pure question	Answers in prose and changes no code
Add to proposal → repos/policycore/core/claims.py + an instruction	File joins the same reviewed diff
Apply a single file	Others stay pending; re-applying is a no-op

- Throughout: `git status` stays clean until Apply. Nothing reaches the repo before then.

05 Cross-team impact & the linked ticket

Needs two logins · Ravi Kumar (1001) + whoever the new ticket is assigned to

1. As Ravi, open AMS-102 or AMS-103 → **Check for other teams affected**.
 2. Confirm a suggested team → a linked Jira ticket is created and assigned.
 3. Log in as that assignee → open their ticket → **Mark as Done**.
 4. Log back in as Ravi.
- After step 2, Stage 1 shows *"1 other team picking up linked work"* and stays **open** — the other team's work is downstream and parallel, so it must not block this ticket.
 - After step 3, the panel clears for Ravi — state comes from the append-only events log, not the browser, which is why it crosses separate logins.
 - No ticket is ever created without a human confirming it.

06 Ad-hoc ticket & the clarity gate

As: Elena Cruz (1002) · **Ticket:** AMS-098 — no linked user story, so it routes to the ad-hoc path

- Runs with **no codebase context** — correctly shows no file-selection panel.
- Three gates share one two-question budget: overall vagueness → a specific missing detail → repo identity.
- A repo match below *high* confidence asks *"is this the right repo?"* rather than scoping against a guess.
- If GitLab is unreachable the repo check is skipped silently and analysis proceeds.

07 Quick-impact chat

Where: "Quick question" box at the top of the S3 page · any login

Try: *"how long would it take to add a beneficiary-change amendment?"*

- Asks at most **two** clarifying questions, then must deliver a final answer.
- Final answer carries impact analysis + hour class + P-level + whether a code change is warranted.
- **Clear** resets the server-side transcript.

08 Problem-record intake — the second user story flavour

Run: `demo/seed_problem_record_ticket.sh` (API must be running)

- A new ticket appears on the board tagged **origin: problem record** with its `PRB...` id.
- It runs the identical downstream flow — only the origin tag differs, which is the point: repeated incidents → problem record → user story converges on the same pipeline.

09 GitLab repo scoping — API only, no UI

```
C=$(mktemp); curl -s -c $C -X POST localhost:8000/api/auth/login \
-H 'Content-Type: application/json' \
-d '{"name":"Ravi Kumar","passcode":"1001"}' >/dev/null

curl -s -b $C -X POST localhost:8000/api/s3/gitlab/scope-auto \
-H 'Content-Type: application/json' \
-d '{"tier_name":"Elite"}' | python -m json.tool
```

- The AI picks a repo from names and descriptions alone, returning confidence, reasoning and alternates.
- Below *high* confidence it returns `needs_clarification` instead of guessing.
- `repo_size` vs `files_reached_llm` shows the funnel economics — only a shortlist is ever fetched, so cost stays flat regardless of repo size.
- Nothing is ever written back to GitLab.

10 Resilience & regression

```
.venv/bin/python -m pytest tests/ -q          # 679 collected
.venv/bin/ruff check .                        # clean
.venv/bin/python tools/verify_s3_live.py --skip-live # offline architecture gate
```

- ☐ Replay works with networking stubbed out entirely.
- ☐ A mid-stream provider failure falls back to replay invisibly.
- ☐ One recording replays correctly for two different tier names.
- ☐ Reset restores baseline in under 10 s.
- ☐ Core recall never drops a required file; screened-out subsystems are never selected.
- ☐ The four checked-in regression suites pass on their own: `polycycore 20 · polycycore_enrolment 19 · claimsportal 5 · enrolldirect 24`. None of those paths appears in any target's `testgen_allowlist` or `codegen_allowlist` — `tests/test_s3_testrun.py` asserts it, and that assertion is the whole value of the beat.

SESSION-DAY RULE

`tools/verify_s3_live.py --gate 10` runs ten real codegen + testgen + pytest cycles from baseline. A cycle passes only if the **live** path succeeded — no silent replay fallback — and the generated tests ran green. **Require ≥ 9/10, otherwise present in LLM_MODE=replay without hesitation.**

11 Onboarding a repo, and the admin panel

As: Manager (9000) · **Where:** `/admin` · **Time:** ~5 min · Answers "what does it take to point this at *our* repo?"

1. Log in as **Manager** → `/admin`. An engineer or tester never sees it — every route depends on `require_manager`, server-side.
 2. **Target apps** — the four services with their live port state; start or stop one without a terminal.
 3. **Onboard a repo** — writes a `.s3targets.json` manifest for a directory under `repos/`.
 4. **Reset environment state** — pick a scope, read its preview, then run it.
- ☐ A source-restoring scope **previews first**: exactly what it restores, what it deletes, and which of those paths are currently dirty. It **409s** rather than overwriting your uncommitted work.
 - ☐ There is **no "reset everything"** scope. Seven explicit ones — `polycycore`, `claimsportal`, `enrolldirect`, `tickets`, `logs`, `proposals`, `caches` — and you name one.
 - ☐ Naming the console as a service is a **400, not an attempt** — it cannot restart the process serving the request.
 - ☐ A manifest written here **needs a console restart** before that target registers. Say so rather than clicking twice.
 - ☐ Drop a `US-YYYY-NNN.md` into `stories/` and reload the board: the ticket is there, unassigned, keyed off the user story id. That is the whole onboarding story — the four steps are in `repos/README.md`.

Known issues

#	ISSUE	IMPACT	STATUS
1	<code>demo/reset_s3.sh</code> and <code>demo/reset_s3_endorsement.sh</code> restore with <code>git checkout HEAD -- repos/...</code> , and the <code>apps/</code> → <code>repos/</code> move was uncommitted, so HEAD did not have those paths	Both PolicyCore resets failed with <i>"pathspec did not match"</i>	Fixed — 3 Aug 2026, by committing the move (<code>e5af8ed</code>); all four resets verified to exit 0. The standing rule: a target move must be committed before these two can restore. The admin panel's <code>reset_blocked_reason</code> check stays, for the next move.
2	US-2026-042 codegen recording stripped ~6 function docstrings from <code>repos/policycore/core/db.py</code>	AMS-102 diff was dominated by unrelated deletions	Fixed — 3 Aug 2026 re-record; docstring counts now identical file for file. Cosmetic blank-line collapsing remains.
3	Subsystem screening was hardcoded to PolicyCore, so the ClaimsPortal target was scored against PolicyCore's decoy subsystems	Panel wrongly reported <code>legacy_platform/settlement</code> 0.49 for AMS-103	Fixed — target-scoped, 2 regression tests
4	Post-apply link always pointed at the PolicyCore portal	AMS-103 sent reviewers to the wrong app	Fixed — per-target app map
5	<code>POST /s3/jira/assign-ticket</code> had no role check at all — assignment was gated only by hiding the button	Any logged-in engineer could assign, and assignment was set-once	Fixed — <code>require_manager</code> server-side; managers can now reassign and unassign
6	AMS-103's Jira card shows <code>APPLICATION: PolicyCore</code> though it is a ClaimsPortal user story	Cosmetic inconsistency in seeded data	Open — low priority
7	Gap gate occasionally asks about a detail the user story already states (e.g. the deductible's data type)	One wasted clarification turn	Open — prompt tuning
8	After QA hand-off the console silently re-targets another ticket, since the handed-off one leaves the developer's filtered board	Disorienting mid-run	Open — low priority
9	ClaimsPortal consoles show no deductible column; <code>payableAmount</code> is not rendered either	No visual before/after on :8081	By design — user story forbids HTML changes

Troubleshooting

SYMPTOM	CAUSE AND FIX
UI looks logged in but every action 401s	Backend restarted (or was run with <code>--reload</code>) and cleared the in-memory session. Log out and back in; restart uvicorn without <code>--reload</code> .

:8081 or :8082 refuses to connect	Neither service is running — start with <code>apps/run-policy-service.sh</code> / <code>apps/run-claims-service.sh</code> (or <code>demo/run_s3_claimsportal.sh</code> for both).
Applied a user story but the ClaimsPortal console is unchanged	Stale process — uvicorn runs without <code>--reload</code> here, so it's still serving the pre-apply code. Ctrl-C and relaunch both services.
A stage will not open	Not a bug — the grey hint under it names the missing precondition (analysis not run, change not applied, ticket not in QA, or you are not the assigned tester).
Console shows results for a ticket you just reset	<code>localStorage</code> cache. Hard-refresh the tab; the reset marker makes the console drop it.
Mutation check returns 409	Generated code drifted from the recording, so no seeded snippet matched. Re-run generate/apply, or re-record. It fails loudly rather than mutating blind.
Analysis is slow on first click	Expected — <code>.cache/llm</code> was cleared by the reset. Run <code>demo/warm_s3_cache.sh</code> to pre-warm.
A PolicyCore reset dies with <code>pathspec 'repos/policycore/app.py' did not match</code>	A target directory was moved and the move is not committed, so those paths are not in HEAD for the script to restore from. Commit the move. Nothing else fixes it.
<code>codegen returned unexpected file set</code>	A repository under <code>repos/</code> was renamed or moved. The committed recordings contain the exact paths, both as keys and inside the generated code's own imports. Restore the directory name — see <code>repos/README.md</code> .
A repo you dropped into <code>repos/</code> does not appear	Either it has no <code>.s3targets.json</code> , or the console has not restarted since you wrote one. A manifest that exists but cannot be parsed <i>raises at import</i> rather than being skipped, so a silent absence means "no manifest", not "bad manifest".
<code>/admin</code> 403s	You are logged in as an engineer or a tester. It is manager-only server-side, not hidden by the UI. Log in as Manager (9000).

All data here is synthetic and the insurer is fictional. No real client data, client names, or credentials appear anywhere in this repo — see `CLAUDE.md` for the full project rules.